

# This is the 2-layer neural network workbook for ECE 239AS Assignment #3

Please follow the notebook linearly to implement a two layer neural network.

Please print out the workbook entirely when completed.

We thank Serena Yeung & Justin Johnson for permission to use code written for the CS 231n class (cs231n.stanford.edu). These are the functions in the cs231n folders and code in the jupyter notebook to preprocess and show the images. The classifiers used are based off of code prepared for CS 231n as well.

The goal of this workbook is to give you experience with training a two layer neural network.

In [1]:

```
import random
import numpy as np
from cs231n.data_utils import load_CIFAR10
import matplotlib.pyplot as plt

%matplotlib inline
%load_ext autoreload
%autoreload 2

def rel_error(x, y):
    """ returns relative error """
    return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))
```

## Toy example

Before loading CIFAR-10, there will be a toy example to test your implementation of the forward and backward pass

In [2]:

```
from nndl.neural_net import TwoLayerNet
```

In [3]:

```
# Create a small net and some toy data to check your implementations.  
# Note that we set the random seed for repeatable experiments.  
  
input_size = 4  
hidden_size = 10  
num_classes = 3  
num_inputs = 5  
  
def init_toy_model():  
    np.random.seed(0)  
    return TwoLayerNet(input_size, hidden_size, num_classes, std=1e-1)  
  
def init_toy_data():  
    np.random.seed(1)  
    X = 10 * np.random.randn(num_inputs, input_size)  
    y = np.array([0, 1, 2, 2, 1])  
    return X, y  
  
net = init_toy_model()  
X, y = init_toy_data()
```

## Compute forward pass scores

In [4]:

```
## Implement the forward pass of the neural network.

# Note, there is a statement if y is None: return scores, which is why
# the following call will calculate the scores.
scores = net.loss(X)
print('Your scores:')
print(scores)
print()
print('correct scores:')
correct_scores = np.asarray([
    [-1.07260209,  0.05083871, -0.87253915],
    [-2.02778743, -0.10832494, -1.52641362],
    [-0.74225908,  0.15259725, -0.39578548],
    [-0.38172726,  0.10835902, -0.17328274],
    [-0.64417314, -0.18886813, -0.41106892]])
print(correct_scores)
print()

# The difference should be very small. We get < 1e-7
print('Difference between your scores and correct scores:')
print(np.sum(np.abs(scores - correct_scores)))
```

Your scores:

```
[[-1.07260209  0.05083871 -0.87253915]
 [-2.02778743 -0.10832494 -1.52641362]
 [-0.74225908  0.15259725 -0.39578548]
 [-0.38172726  0.10835902 -0.17328274]
 [-0.64417314 -0.18886813 -0.41106892]]
```

correct scores:

```
[[-1.07260209  0.05083871 -0.87253915]
 [-2.02778743 -0.10832494 -1.52641362]
 [-0.74225908  0.15259725 -0.39578548]
 [-0.38172726  0.10835902 -0.17328274]
 [-0.64417314 -0.18886813 -0.41106892]]
```

Difference between your scores and correct scores:

```
3.381231233889892e-08
```

## Forward pass loss

In [5]:

```
loss, _ = net.loss(X, y, reg=0.05)
correct_loss = 1.071696123862817

# should be very small, we get < 1e-12
print("Loss:", loss)
print('Difference between your loss and correct loss:')
print(np.sum(np.abs(loss - correct_loss)))
```

```
Loss: 1.071696123862817
Difference between your loss and correct loss:
0.0
```

## Backward pass

Implements the backwards pass of the neural network. Check your gradients with the gradient check utilities provided.

In [6]:

```
from cs231n.gradient_check import eval_numerical_gradient

# Use numeric gradient checking to check your implementation of the backward pass.
# If your implementation is correct, the difference between the numeric and
# analytic gradients should be less than 1e-8 for each of W1, W2, b1, and b2.

loss, grads = net.loss(X, y, reg=0.05)

# these should all be less than 1e-8 or so
for param_name in grads:
    f = lambda W: net.loss(X, y, reg=0.05)[0]
    param_grad_num = eval_numerical_gradient(f, net.params[param_name], verbose=False)
    print('{} max relative error: {}'.format(param_name, rel_error(param_grad_num, grads[param_name])))
```

```
W1 max relative error: 1.2832823337649917e-09
W2 max relative error: 2.9632227682005116e-10
b1 max relative error: 3.172680092703762e-09
b2 max relative error: 1.248270530283678e-09
```

## Training the network

Implement `neural_net.train()` to train the network via stochastic gradient descent, much like the softmax and SVM.

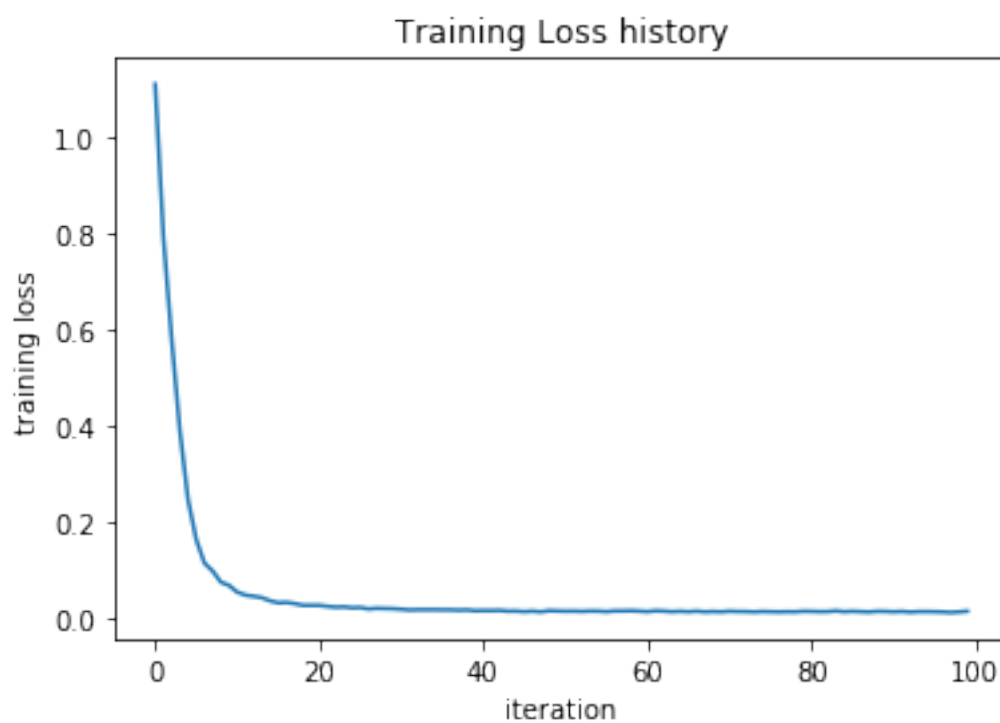
In [7]:

```
net = init_toy_model()
stats = net.train(X, y, X, y,
                  learning_rate=1e-1, reg=5e-6,
                  num_iters=100, verbose=False)

print('Final training loss: ', stats['loss_history'][-1])

# plot the loss history
plt.plot(stats['loss_history'])
plt.xlabel('iteration')
plt.ylabel('training loss')
plt.title('Training Loss history')
plt.show()
```

Final training loss: 0.014497864587765886



## Classify CIFAR-10

Do classification on the CIFAR-10 dataset.

In [8]:

```
from cs231n.data_utils import load_CIFAR10

def get_CIFAR10_data(num_training=49000, num_validation=1000, num_test=1000):
    """
    Load the CIFAR-10 dataset from disk and perform preprocessing to prepare
    it for the two-layer neural net classifier. These are the same steps as
    we used for the SVM, but condensed to a single function.
    """
    # Load the raw CIFAR-10 data
    cifar10_dir = '/Users/egecetintas/Desktop/UCLA/c247/hw2/cifar-10-batches-p
y' # You need to update this line
    X_train, y_train, X_test, y_test = load_CIFAR10(cifar10_dir)

    # Subsample the data
    mask = list(range(num_training, num_training + num_validation))
    X_val = X_train[mask]
    y_val = y_train[mask]
    mask = list(range(num_training))
    X_train = X_train[mask]
    y_train = y_train[mask]
    mask = list(range(num_test))
    X_test = X_test[mask]
    y_test = y_test[mask]

    # Normalize the data: subtract the mean image
    mean_image = np.mean(X_train, axis=0)
    X_train -= mean_image
    X_val -= mean_image
    X_test -= mean_image

    # Reshape data to rows
    X_train = X_train.reshape(num_training, -1)
    X_val = X_val.reshape(num_validation, -1)
    X_test = X_test.reshape(num_test, -1)

    return X_train, y_train, X_val, y_val, X_test, y_test

# Invoke the above function to get our data.
X_train, y_train, X_val, y_val, X_test, y_test = get_CIFAR10_data()
print('Train data shape: ', X_train.shape)
print('Train labels shape: ', y_train.shape)
print('Validation data shape: ', X_val.shape)
print('Validation labels shape: ', y_val.shape)
print('Test data shape: ', X_test.shape)
print('Test labels shape: ', y_test.shape)
```

```
Train data shape: (49000, 3072)
Train labels shape: (49000,)
Validation data shape: (1000, 3072)
Validation labels shape: (1000,)
Test data shape: (1000, 3072)
Test labels shape: (1000,)
```

## Running SGD

If your implementation is correct, you should see a validation accuracy of around 28-29%.

In [9]:

```
input_size = 32 * 32 * 3
hidden_size = 50
num_classes = 10
net = TwoLayerNet(input_size, hidden_size, num_classes)
# Train the network
stats = net.train(X_train, y_train, X_val, y_val,
                  num_iters=1000, batch_size=200,
                  learning_rate=1e-4, learning_rate_decay=0.95,
                  reg=0.25, verbose=True)

# Predict on the validation set
val_acc = (net.predict(X_val) == y_val).mean()
print('Validation accuracy: ', val_acc)

# Save this net as the variable subopt_net for later comparison.
subopt_net = net
```

```
iteration 0 / 1000: loss 2.302757518613176
iteration 100 / 1000: loss 2.302120159207236
iteration 200 / 1000: loss 2.2956136007408703
iteration 300 / 1000: loss 2.2518259043164135
iteration 400 / 1000: loss 2.1889952350467756
iteration 500 / 1000: loss 2.1162527791897747
iteration 600 / 1000: loss 2.064670827698217
iteration 700 / 1000: loss 1.990168862308394
iteration 800 / 1000: loss 2.0028276401246856
iteration 900 / 1000: loss 1.9465176817856495
Validation accuracy: 0.283
```

## Questions:

The training accuracy isn't great.

(1) What are some of the reasons why this is the case? Take the following cell to do some analyses and then report your answers in the cell following the one below.

(2) How should you fix the problems you identified in (1)?

In [10]:

```
stats['train_acc_history']
```

Out[10]:

```
[0.095, 0.15, 0.25, 0.25, 0.315]
```

In [11]:

```

# ===== #

# YOUR CODE HERE:
# Do some debugging to gain some insight into why the optimization
# isn't great.
# ===== #

# Plot the loss function and train / validation accuracies
fig = plt.figure(figsize=(2*6.4, 4.8))
plt.subplot(121)
plt.plot(stats['train_acc_history'],label='train_acc_history')
plt.plot(stats['val_acc_history'],label='val_acc_history')
plt.title('Hyperparameters: \nnum_iters=1000, batch_size=200, learning_rate=1e-4, learning_rate_decay=0.95, reg=0.25')
plt.legend()
plt.subplot(122)
plt.title('Loss History')
plt.plot(stats['loss_history'])
plt.show()

#Trying different Number of Iteration
#1
net = TwoLayerNet(input_size, hidden_size, num_classes)
stats = net.train(X_train, y_train, X_val, y_val,
                  num_iters=10000, batch_size=200,
                  learning_rate=1e-4, learning_rate_decay=0.95,
                  reg=0.25, verbose=False)

val_acc = (net.predict(X_val) == y_val).mean()
print('Validation accuracy: ', val_acc)
fig = plt.figure(figsize=(2*6.4, 4.8))
plt.subplot(121)
plt.plot(stats['train_acc_history'],label='train_acc_history')
plt.plot(stats['val_acc_history'],label='val_acc_history')
plt.title('Hyperparameters: \nnum_iters=10000, batch_size=200, learning_rate=1e-4, learning_rate_decay=0.95, reg=0.25')
plt.legend()
plt.subplot(122)
plt.title('Loss History')
plt.plot(stats['loss_history'])
plt.show()

#2
net = TwoLayerNet(input_size, hidden_size, num_classes)
stats = net.train(X_train, y_train, X_val, y_val,
                  num_iters=10000, batch_size=200,
                  learning_rate=1e-3, learning_rate_decay=0.95,
                  reg=0.25, verbose=False)
val_acc = (net.predict(X_val) == y_val).mean()
print('Validation accuracy: ', val_acc)
fig = plt.figure(figsize=(2*6.4, 4.8))
plt.subplot(121)
plt.plot(stats['train_acc_history'],label='train_acc_history')
plt.plot(stats['val_acc_history'],label='val_acc_history')
plt.title('Hyperparameters: \nnum_iters=10000, batch_size=200, learning_rate=1e-3, learning_rate_decay=0.95, reg=0.25')
plt.legend()
plt.subplot(122)

```



```

plt.title('Loss History')

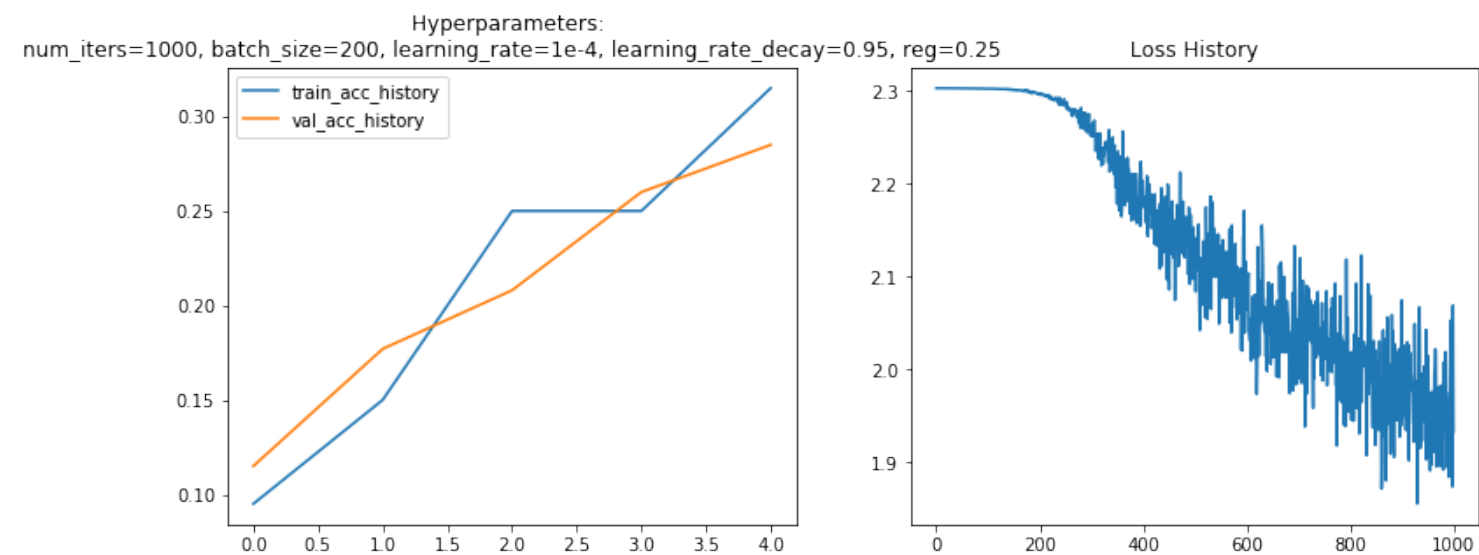
plt.plot(stats['loss_history'])
plt.show()

#3
net = TwoLayerNet(input_size, hidden_size, num_classes)
stats = net.train(X_train, y_train, X_val, y_val,
                  num_iters=10000, batch_size=200,
                  learning_rate=1e-3, learning_rate_decay=0.95,
                  reg=0.05, verbose=False)
val_acc = (net.predict(X_val) == y_val).mean()
print('Validation accuracy: ', val_acc)
fig = plt.figure(figsize=(2*6.4, 4.8))
plt.subplot(121)
plt.plot(stats['train_acc_history'], label='train_acc_history')
plt.plot(stats['val_acc_history'], label='val_acc_history')
plt.title('Hyperparameters: \nnum_iters=10000, batch_size=200, learning_rate=1e-3, learning_rate_decay=0.95, reg=0.05')
plt.legend()
plt.subplot(122)
plt.title('Loss History')
plt.plot(stats['loss_history'])
plt.show()

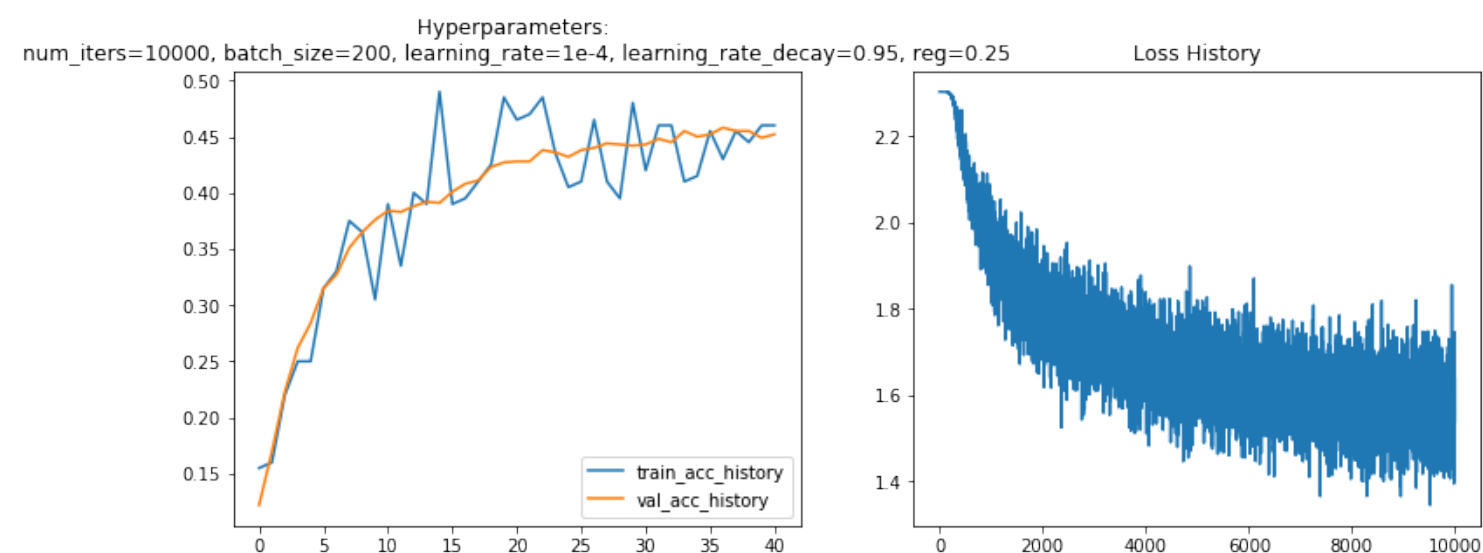
pass

# ===== #
# END YOUR CODE HERE
# ===== #

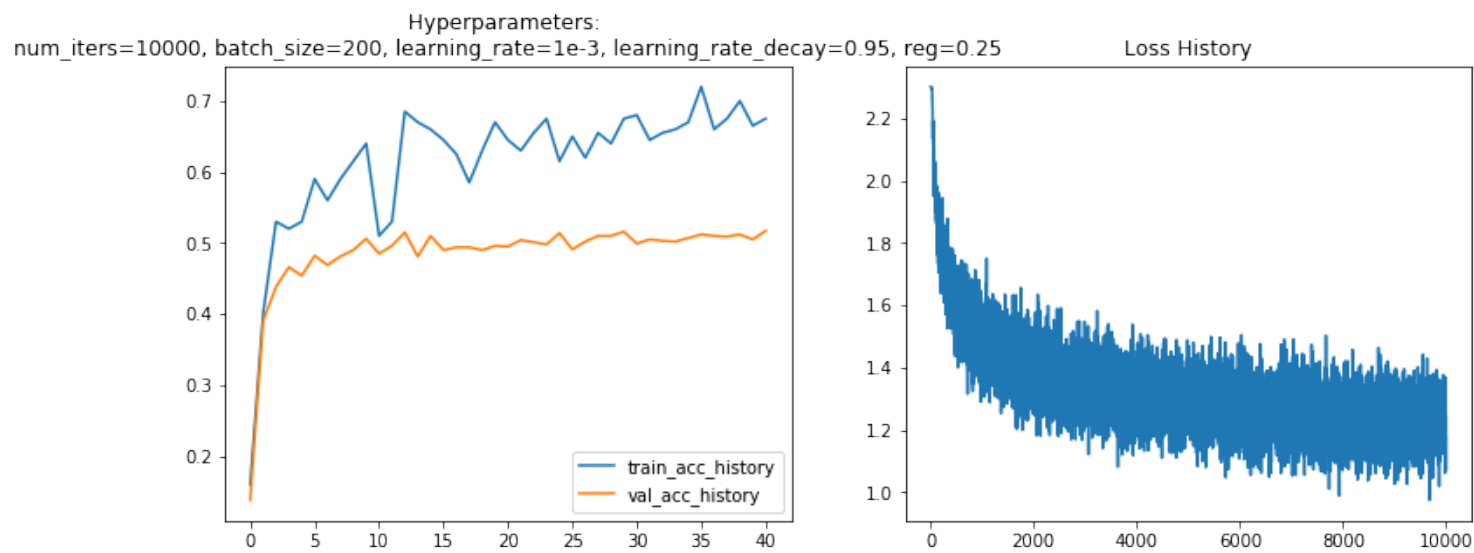
```



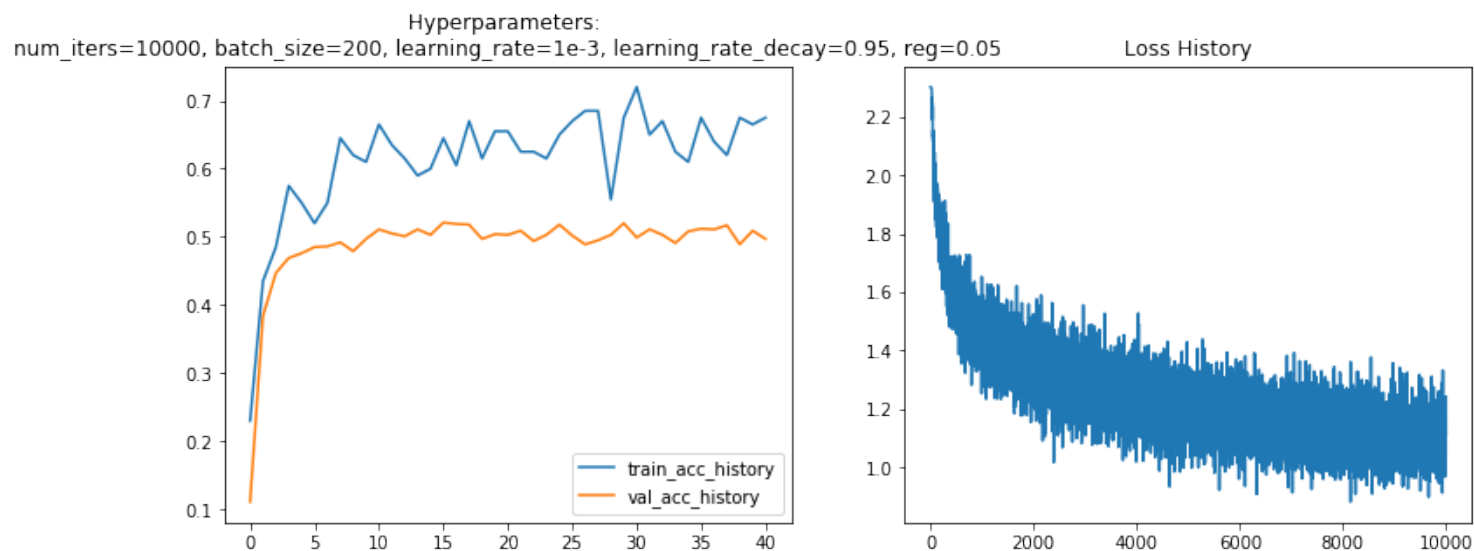
Validation accuracy: 0.45



Validation accuracy: 0.514



Validation accuracy: 0.499



## Answers:

(1) The plots of loss and accuracy history above suggests that we don't achieve adequate fitting of the training data yet. Training accuracy and validation accuracy follows a similar trend -both increasing linearly- and it would suggest that there is still room for more learning. If we have a look at the loss history, we can observe that the decrease trend is still pretty linear for 1000 iteration suggesting that SGD does not find a local minima yet.

(2) In order to fix the problem, we can increase the number of iterations to force the model to learn the training data even more. Alternatively, we can play with some of the hyperparameters such as learning rate or regularization strength to optimize for a specific number of iterations.

## Optimize the neural network

Use the following part of the Jupyter notebook to optimize your hyperparameters on the validation set. Store your nets as `best_net`.

In [12]:

```
best_net = None # store the best model into this

# ===== #
# YOUR CODE HERE:
```

```

# YOUR CODE HERE:
# Optimize over your hyperparameters to arrive at the best neural
# network. You should be able to get over 50% validation accuracy.
# For this part of the notebook, we will give credit based on the
# accuracy you get. Your score on this question will be multiplied by:
# min(floor((X - 28%) / %22, 1)
# where if you get 50% or higher validation accuracy, you get full
# points.
#
# Note, you need to use the same network structure (keep hidden_size = 50)!
# ===== #
input_size = 32 * 32 * 3
hidden_size = 50
num_classes = 10

optimLength = 200

val_acc_list = []
learning_rates = []
batch_list = []
iter_list = []
reg_list = []

start_reg = 0.15
end_reg = 0.5
start_iter = 100
end_iter = 100
start_lr = -7
end_lr = -3.5
start_batch = 100
end_batch = 100
net = TwoLayerNet(input_size, hidden_size, num_classes)
for i in range(optimLength):
    #Generate Random Hyperparameters
    learning_rate_random = 10 ** np.random.uniform(start_lr, end_lr,1)
    num_iters_random = np.random.uniform(start_iter, end_iter,1)
    reg_random = np.random.uniform(start_reg,end_reg,1)
    batch_random = int(np.random.uniform(start_batch,end_batch,1))
    learning_rates.append(learning_rate_random)
    reg_list.append(reg_random)
    iter_list.append(num_iters_random)
    batch_list.append(batch_random)

    stats = net.train(X_train, y_train, X_val, y_val,
                      num_iters=num_iters_random, batch_size=batch_random,
                      learning_rate=learning_rate_random, learning_rate_decay=0.95,
                      reg=reg_random, verbose=False)
    val_acc = (net.predict(X_val) == y_val).mean()
    print('Validation accuracy: ', val_acc)
    val_acc_list.append(val_acc)

    hyperparamers = sorted(zip(val_acc_list,learning_rates,iter_list,reg_list,
batch_list))

    k = 10 # Intervals for Random Hyperparameters are updated every k training
.
    if (i+1) % k == 0:
        #Update Random Hyperparameters' Generation Interval

```

```

print('\n')

print('Current Epoch: ', str((i+1)/k))
print('max_val_acc:', max(hyperparameters[-k:])[0])
print('learning_rate:', max(hyperparameters[-k:])[1])
print('iter_num:', max(hyperparameters[-k:])[2])
print('reg_num:', max(hyperparameters[-k:])[3])
print('batch_num:', max(hyperparameters[-k:])[4])
print('Hyperparameter limits are updated for the next epoch!')
print('\n')
start_lr = np.log10(max(hyperparameters[-k:])[1]) - 0.75*0.95
end_lr = np.log10(max(hyperparameters[-k:])[1]) + 0.75*0.95
start_reg = max(hyperparameters[-k:])[3] - 0.05*0.95
end_reg = max(hyperparameters[-k:])[3] + 0.1*0.95

```

```

stats = net.train(X_train, y_train, X_val, y_val,
                  num_iters=100, batch_size=100,
                  learning_rate=max(hyperparameters[-5:])[1][0], learning_rate_decay=0.95,
                  reg=max(hyperparameters[-5:])[3][0], verbose=False)
best_net = net

```

```

pass
#Plot the val_acc history for every set of hyperparameters
plt.plot(val_acc_list)
plt.xlabel('Index of Different Hyperparameters')
plt.ylabel('Validation Loss')
# ===== #
# END YOUR CODE HERE
# ===== #
val_acc = (best_net.predict(X_val) == y_val).mean()
print('Validation accuracy: ', val_acc)

```

```

Validation accuracy: 0.17
Validation accuracy: 0.198
Validation accuracy: 0.196
Validation accuracy: 0.196
Validation accuracy: 0.196
Validation accuracy: 0.195
Validation accuracy: 0.195
Validation accuracy: 0.195
Validation accuracy: 0.194
Validation accuracy: 0.195

```

```

Current Epoch: 1.0
max_val_acc: 0.198
learning_rate: [0.00012854]
iter_num: [100.]
reg_num: [0.30677967]
batch_num: 100
Hyperparameter limits are updated for the next epoch!

```

```

Validation accuracy: 0.175
Validation accuracy: 0.222

```

Validation accuracy: 0.247  
Validation accuracy: 0.258  
Validation accuracy: 0.301  
Validation accuracy: 0.32  
Validation accuracy: 0.334  
Validation accuracy: 0.338  
Validation accuracy: 0.366  
Validation accuracy: 0.368

Current Epoch: 2.0  
max\_val\_acc: 0.368  
learning\_rate: [3.02470654e-05]  
iter\_num: [100.]  
reg\_num: [0.31283904]  
batch\_num: 100  
Hyperparameter limits are updated for the next epoch!

Validation accuracy: 0.366  
Validation accuracy: 0.363  
Validation accuracy: 0.364  
Validation accuracy: 0.366  
Validation accuracy: 0.371  
Validation accuracy: 0.376  
Validation accuracy: 0.373  
Validation accuracy: 0.372  
Validation accuracy: 0.385  
Validation accuracy: 0.382

Current Epoch: 3.0  
max\_val\_acc: 0.385  
learning\_rate: [7.13372609e-05]  
iter\_num: [100.]  
reg\_num: [0.33159639]  
batch\_num: 100  
Hyperparameter limits are updated for the next epoch!

Validation accuracy: 0.381  
Validation accuracy: 0.392  
Validation accuracy: 0.39  
Validation accuracy: 0.398  
Validation accuracy: 0.4  
Validation accuracy: 0.404  
Validation accuracy: 0.408  
Validation accuracy: 0.408  
Validation accuracy: 0.427  
Validation accuracy: 0.424

Current Epoch: 4.0  
max\_val\_acc: 0.427  
learning\_rate: [0.00033178]  
iter\_num: [100.]  
reg\_num: [0.32871068]

batch\_num: 100  
Hyperparameter limits are updated for the next epoch!

Validation accuracy: 0.433  
Validation accuracy: 0.421  
Validation accuracy: 0.44  
Validation accuracy: 0.427  
Validation accuracy: 0.465  
Validation accuracy: 0.452  
Validation accuracy: 0.458  
Validation accuracy: 0.414  
Validation accuracy: 0.455  
Validation accuracy: 0.45

Current Epoch: 5.0  
max\_val\_acc: 0.465  
learning\_rate: [0.00018463]  
iter\_num: [100.]  
reg\_num: [0.3408659]  
batch\_num: 100  
Hyperparameter limits are updated for the next epoch!

Validation accuracy: 0.461  
Validation accuracy: 0.46  
Validation accuracy: 0.469  
Validation accuracy: 0.474  
Validation accuracy: 0.477  
Validation accuracy: 0.465  
Validation accuracy: 0.456  
Validation accuracy: 0.482  
Validation accuracy: 0.482  
Validation accuracy: 0.474

Current Epoch: 6.0  
max\_val\_acc: 0.482  
learning\_rate: [0.00015447]  
iter\_num: [100.]  
reg\_num: [0.43207708]  
batch\_num: 100  
Hyperparameter limits are updated for the next epoch!

Validation accuracy: 0.494  
Validation accuracy: 0.492  
Validation accuracy: 0.49  
Validation accuracy: 0.5  
Validation accuracy: 0.477  
Validation accuracy: 0.494  
Validation accuracy: 0.501  
Validation accuracy: 0.499  
Validation accuracy: 0.496  
Validation accuracy: 0.497

Current Epoch: 7.0  
max\_val\_acc: 0.501  
learning\_rate: [5.57046218e-05]  
iter\_num: [100.]  
reg\_num: [0.42690496]  
batch\_num: 100  
Hyperparameter limits are updated for the next epoch!

Validation accuracy: 0.499  
Validation accuracy: 0.504  
Validation accuracy: 0.502  
Validation accuracy: 0.497  
Validation accuracy: 0.495  
Validation accuracy: 0.5  
Validation accuracy: 0.501  
Validation accuracy: 0.497  
Validation accuracy: 0.499  
Validation accuracy: 0.501

Current Epoch: 8.0  
max\_val\_acc: 0.504  
learning\_rate: [4.81385873e-05]  
iter\_num: [100.]  
reg\_num: [0.49469929]  
batch\_num: 100  
Hyperparameter limits are updated for the next epoch!

Validation accuracy: 0.49  
Validation accuracy: 0.489  
Validation accuracy: 0.504  
Validation accuracy: 0.491  
Validation accuracy: 0.495  
Validation accuracy: 0.495  
Validation accuracy: 0.494  
Validation accuracy: 0.499  
Validation accuracy: 0.496  
Validation accuracy: 0.498

Current Epoch: 9.0  
max\_val\_acc: 0.504  
learning\_rate: [5.35478816e-05]  
iter\_num: [100.]  
reg\_num: [0.48770629]  
batch\_num: 100  
Hyperparameter limits are updated for the next epoch!

Validation accuracy: 0.497  
Validation accuracy: 0.494  
Validation accuracy: 0.509  
Validation accuracy: 0.505  
Validation accuracy: 0.498

Validation accuracy: 0.502  
Validation accuracy: 0.494  
Validation accuracy: 0.491  
Validation accuracy: 0.484  
Validation accuracy: 0.495

Current Epoch: 10.0

max\_val\_acc: 0.509

learning\_rate: [0.00010335]

iter\_num: [100.]

reg\_num: [0.479229]

batch\_num: 100

Hyperparameter limits are updated for the next epoch!

Validation accuracy: 0.493  
Validation accuracy: 0.492  
Validation accuracy: 0.501  
Validation accuracy: 0.506  
Validation accuracy: 0.503  
Validation accuracy: 0.496  
Validation accuracy: 0.503  
Validation accuracy: 0.506  
Validation accuracy: 0.489  
Validation accuracy: 0.506

Current Epoch: 11.0

max\_val\_acc: 0.509

learning\_rate: [0.00010335]

iter\_num: [100.]

reg\_num: [0.479229]

batch\_num: 100

Hyperparameter limits are updated for the next epoch!

Validation accuracy: 0.495  
Validation accuracy: 0.507  
Validation accuracy: 0.497  
Validation accuracy: 0.504  
Validation accuracy: 0.512  
Validation accuracy: 0.497  
Validation accuracy: 0.505  
Validation accuracy: 0.503  
Validation accuracy: 0.509  
Validation accuracy: 0.516

Current Epoch: 12.0

max\_val\_acc: 0.516

learning\_rate: [0.00034629]

iter\_num: [100.]

reg\_num: [0.53231348]

batch\_num: 100

Hyperparameter limits are updated for the next epoch!



Validation accuracy: 0.501  
Validation accuracy: 0.484  
Validation accuracy: 0.498  
Validation accuracy: 0.501  
Validation accuracy: 0.412  
Validation accuracy: 0.486  
Validation accuracy: 0.476  
Validation accuracy: 0.499  
Validation accuracy: 0.514  
Validation accuracy: 0.463

Current Epoch: 13.0  
max\_val\_acc: 0.516  
learning\_rate: [0.00034629]  
iter\_num: [100.]  
reg\_num: [0.53231348]  
batch\_num: 100  
Hyperparameter limits are updated for the next epoch!

Validation accuracy: 0.48  
Validation accuracy: 0.465  
Validation accuracy: 0.495  
Validation accuracy: 0.476  
Validation accuracy: 0.497  
Validation accuracy: 0.435  
Validation accuracy: 0.448  
Validation accuracy: 0.468  
Validation accuracy: 0.502  
Validation accuracy: 0.5

Current Epoch: 14.0  
max\_val\_acc: 0.516  
learning\_rate: [0.00034629]  
iter\_num: [100.]  
reg\_num: [0.53231348]  
batch\_num: 100  
Hyperparameter limits are updated for the next epoch!

Validation accuracy: 0.481  
Validation accuracy: 0.508  
Validation accuracy: 0.51  
Validation accuracy: 0.511  
Validation accuracy: 0.499  
Validation accuracy: 0.459  
Validation accuracy: 0.507  
Validation accuracy: 0.441  
Validation accuracy: 0.479  
Validation accuracy: 0.498

Current Epoch: 15.0  
max\_val\_acc: 0.516

learning\_rate: [0.00034629]  
iter\_num: [100.]  
reg\_num: [0.53231348]  
batch\_num: 100  
Hyperparameter limits are updated for the next epoch!

Validation accuracy: 0.497  
Validation accuracy: 0.517  
Validation accuracy: 0.492  
Validation accuracy: 0.513  
Validation accuracy: 0.506  
Validation accuracy: 0.508  
Validation accuracy: 0.522  
Validation accuracy: 0.515  
Validation accuracy: 0.526  
Validation accuracy: 0.524

Current Epoch: 16.0  
max\_val\_acc: 0.526  
learning\_rate: [0.00010027]  
iter\_num: [100.]  
reg\_num: [0.5116973]  
batch\_num: 100  
Hyperparameter limits are updated for the next epoch!

Validation accuracy: 0.512  
Validation accuracy: 0.521  
Validation accuracy: 0.514  
Validation accuracy: 0.518  
Validation accuracy: 0.525  
Validation accuracy: 0.498  
Validation accuracy: 0.508  
Validation accuracy: 0.52  
Validation accuracy: 0.525  
Validation accuracy: 0.522

Current Epoch: 17.0  
max\_val\_acc: 0.526  
learning\_rate: [0.00010027]  
iter\_num: [100.]  
reg\_num: [0.5116973]  
batch\_num: 100  
Hyperparameter limits are updated for the next epoch!

Validation accuracy: 0.523  
Validation accuracy: 0.511  
Validation accuracy: 0.519  
Validation accuracy: 0.528  
Validation accuracy: 0.494  
Validation accuracy: 0.52  
Validation accuracy: 0.527  
Validation accuracy: 0.506

Validation accuracy: 0.514  
Validation accuracy: 0.518

Current Epoch: 18.0  
max\_val\_acc: 0.528  
learning\_rate: [6.5681754e-05]  
iter\_num: [100.]  
reg\_num: [0.50103815]  
batch\_num: 100  
Hyperparameter limits are updated for the next epoch!

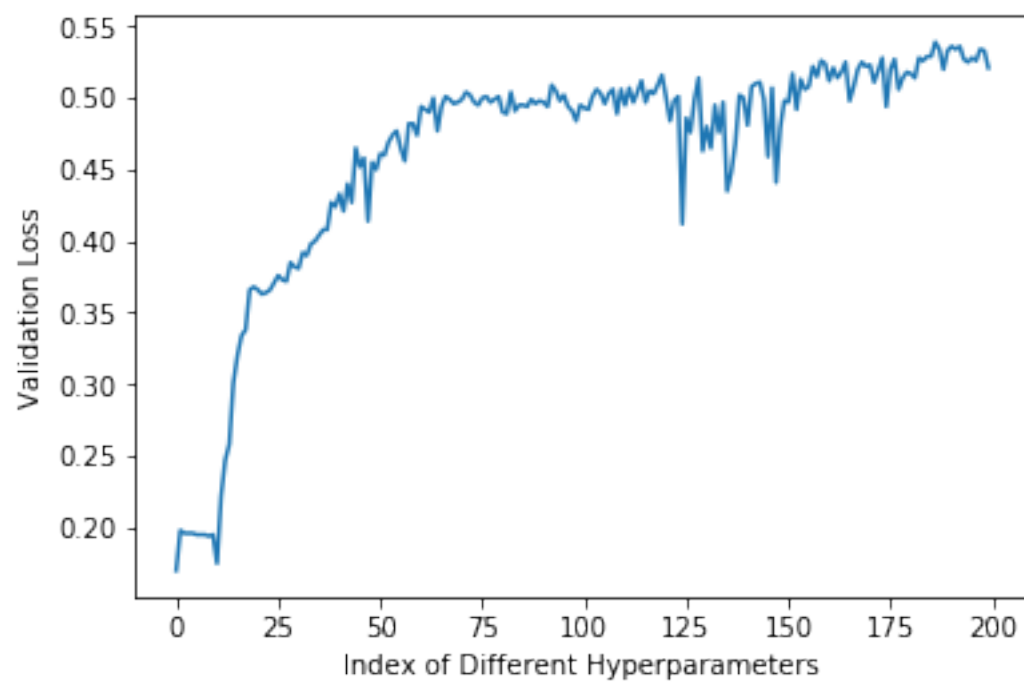
Validation accuracy: 0.517  
Validation accuracy: 0.514  
Validation accuracy: 0.528  
Validation accuracy: 0.526  
Validation accuracy: 0.529  
Validation accuracy: 0.529  
Validation accuracy: 0.539  
Validation accuracy: 0.534  
Validation accuracy: 0.52  
Validation accuracy: 0.533

Current Epoch: 19.0  
max\_val\_acc: 0.539  
learning\_rate: [2.11191102e-05]  
iter\_num: [100.]  
reg\_num: [0.52502032]  
batch\_num: 100  
Hyperparameter limits are updated for the next epoch!

Validation accuracy: 0.536  
Validation accuracy: 0.534  
Validation accuracy: 0.536  
Validation accuracy: 0.527  
Validation accuracy: 0.525  
Validation accuracy: 0.528  
Validation accuracy: 0.526  
Validation accuracy: 0.534  
Validation accuracy: 0.533  
Validation accuracy: 0.521

Current Epoch: 20.0  
max\_val\_acc: 0.539  
learning\_rate: [2.11191102e-05]  
iter\_num: [100.]  
reg\_num: [0.52502032]  
batch\_num: 100  
Hyperparameter limits are updated for the next epoch!

Validation accuracy: 0.536



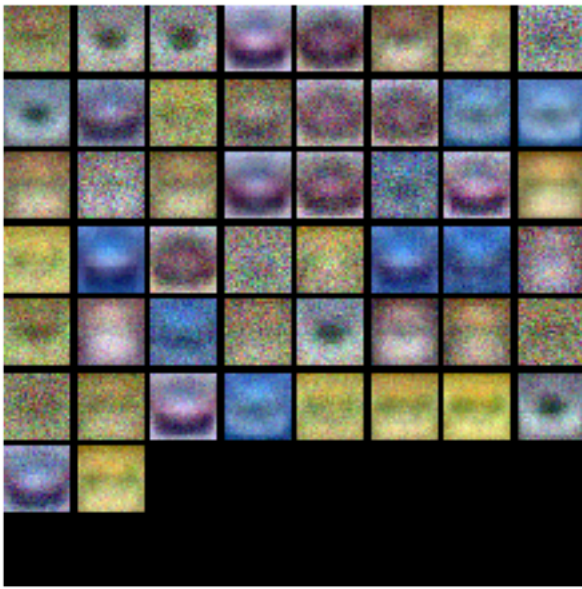
In [13]:

```
from cs231n.vis_utils import visualize_grid

# Visualize the weights of the network

def show_net_weights(net):
    W1 = net.params['W1']
    W1 = W1.T.reshape(32, 32, 3, -1).transpose(3, 0, 1, 2)
    plt.imshow(visualize_grid(W1, padding=3).astype('uint8'))
    plt.gca().axis('off')
    plt.show()

show_net_weights(subopt_net)
show_net_weights(best_net)
```



## Question:

(1) What differences do you see in the weights between the suboptimal net and the best net you arrived at?

## Answer:

(1) The weights of the best net are sharper than suboptimal net. Suboptimal net's weights are more blurred. For instance, we can reach to this conclusion by looking at the car shaped weights. The shapes are more visible in the best network's weights comparing to suboptimal net's weights.

## Evaluate on test set

In [14]:

```
test_acc = (best_net.predict(X_test) == y_test).mean()
print('Test accuracy: ', test_acc)
```

Test accuracy: 0.525

In [ ]: